

Files



Character Operations

- ❖ Comparing a character `c` and an **element** in a string `s1`

```
if ( s1[i] == 'c' ) ...
```

Single quotes

- ❖ Changing a single character in an array

```
s1[i] = 'c';
```

Note: You do not need the **single** quotes if the character is a variable.

```
char letter;  
char str[10];  
letter = 'm';  
str[3] = letter;
```

ASCII

- ❖ **ASCII** - American Standard Code for Information Interchange
- ❖ International counterpart is ISO646.
- ❖ It is a 7-bit code (128 values).
- ❖ Uses one byte (8 bits) to store a character.

ASCII

- ❖ Upper and lower case differ by one bit.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A

- ❖ First bit is always 0 in pure ASCII.

<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
.	46		2E

ASCII

- ❖ Upper and lower case differ by one bit.
- ❖ First bit is always 0 in pure ASCII.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A
<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
.	46		2E

'a' is a 1 byte number

We think of it as *a* ,
C thinks of it as 97

ASCII

- ❖ Upper and lower case differ by one bit.
- ❖ First bit is always 0 in pure ASCII.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A
<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
	46		2E

'a' is a 1 byte number

We think of it as *a* ,
C thinks of it as 97

'a' < 'b' same as
97 < 98 true

ASCII

- ❖ Upper and lower case differ by one bit.
- ❖ First bit is always 0 in pure ASCII.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A
<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
	46		2E

'a' is a 1 byte number

We think of it as *a* ,
C thinks of it as 97

'a' > 'A' same as
97 > 65 true

ASCII

- ❖ Upper and lower case differ by one bit.
- ❖ First bit is always 0 in pure ASCII.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A
<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
	46		2E

'a' is a 1 byte number

We think of it as *a* ,
C thinks of it as 97

'C' + ('a' - 'A') == 'c' same as
67 + 32 == 99 true

ASCII

- ❖ Upper and lower case differ by one bit.
- ❖ First bit is always 0 in pure ASCII.

Character	Decimal	Bits	Hex
A	65	0100 0001	41
Z	90	0101 1010	5A
a	97	0110 0001	61
z	122	0111 1010	7A
<space>	32		20
<line feed \n>	10		0A
<tab \t>	9		09
0	48		30
9	57		39
	46		2E

'a' is a 1 byte number

We think of it as *a* ,
C thinks of it as 97

```
char c = (char) 97;  
c == 'a'    true
```

Unicode

- ❖ Other character encodings exist for characters that do not appear in the ASCII set, *i.e.* international character sets.
- ❖ **Unicode** Text Format (UTF) is supported by Java (and other languages) and uses between 1 and 4 bytes to store a character. An ASCII character is a 1-byte UTF character .
- ❖ *From Wikipedia:* Unicode contains 137,994 characters (consisting of 137,766 graphic characters, 163 format characters and 65 control characters) covering 150 modern and historic scripts, as well as multiple symbol sets and emoji.

Text File

- ❖ A text file is a **sequence** of **ASCII** characters stored on a disk.
- ❖ Each character is stored as **one** byte.
- ❖ Similar to strings in C.

Text File Format Example

- ❖ Store the string “cost 2.98” in a text file. Each character is 1 byte.
- ❖ When the text file is read every byte is read into a char.

c	o	s	t		2	.	9	8
01100011	01101111	01110011	01110100	00100000	00110010	00101110	00111001	00111000

- ❖ Note: numbers are stored as character strings so 2.98 needs 4 bytes or characters - it is not stored as a float! The integer 2 would be stored in 1 byte.

Text Files

- ❖ Text editors are simply applications that read / display / write text files.
- ❖ The only difference between **text** and **binary** files is the type of data stored in the files.
- ❖ **Text** files contain only **ASCII** characters.
- ❖ Both types are read from and written to using the same techniques.

Binary Files

- ❖ **Binary** files contain **non-ASCII** characters.
- ❖ **Text** is still stored using **ASCII**
- ❖ **Numbers** are stored in their **binary** format
- ❖ **Structures** can also be stored

Binary Files

- ❖ Store the integer number 10383 in a binary file.
- ❖ An integer is stored as 4 bytes in a binary format.
 $10383 = 00000000\ 00000000\ 00010100\ 10001111$
- ❖ If it was stored in a text file it would occupy 5 bytes.
- ❖ Binary representations can take less space but cannot be edited by a text editor.

Interpreting Data in Files

- ❖ Text and binary data can be mixed in the same file.
- ❖ Unix / Linux does not care what is in the file, it is all dependant on how it is interpreted by a program.
- ❖ When reading a file you must read it in the same order in which the file was written or the data will be corrupted.
- ❖ The system considers the file to be a **collection of bytes**.

0000 0000 0001 0100 1000 1111

- ❖ These 4 bytes could be interpreted as 4 characters (1 byte each) or an integer (4 bytes).
- ❖ But if we interpret this as 4 characters - nonsense!
- ❖ As an integer it is: 5263

	Binary	Hex	ASCII
1st Byte	0000 0000	0	NUL
2nd Byte	0000 0000	0	NUL
3rd Byte	0001 0100	14	DC4
4th Byte	1000 1111	8F	undefined (> 127)

0000 0000 0001 0100 1000 1111

- ❖ In a C program:

```
fscanf ( fptr, "%d", &i );
```

```
fscanf ( fptr, "%c%c%c%c", &a, &b, &c, &d );
```

- ❖ If this bit string (4 bytes worth) was in a file than it could not be read in a text editor since some of the characters are not ASCII.
- ❖ See program: *interpretBytes.c*


```
#include <stdio.h>
```

```
int main ( )
```

```
{
```

```
    unsigned int num = 5263;
```

```
    char a, b, c, d;
```

```
    FILE *fptr;
```

```
    fptr = fopen ( "testme", "wb" );
```

```
    fwrite ( &num, sizeof(num),1, fptr );
```

```
    fclose ( fptr );
```

```
    fptr = fopen ( "testme", "r" );
```

```
    fscanf ( fptr, "%d", &num );
```

```
    printf ( "Number = %d\n", num );
```

```
    fclose ( fptr );
```

```
    fptr = fopen ( "testme", "r" );
```

```
    fscanf ( fptr, "%c%c%c%c", &a, &b, &c, &d );
```

```
    printf ( "a = %c \n", a );
```

```
    printf ( "b = %c \n", b );
```

```
    printf ( "c = %c \n", c );
```

```
    printf ( "d = %c \n", d );
```

```
    fclose ( fptr );
```

```
    return(0);
```

```
}
```

```
$ ./interpretBytes
```

```
Number = 5263
```

```
a = ?
```

```
b =
```

```
c =
```

```
d =
```

```
$ od -d testme
```

```
00000000      5263      0
```

```
00000004
```

```
$ od -c testme
```

```
00000000  217 024  \0  \0
```

```
00000004
```

```
$
```

File Operations

- ❖ Files in C are treated as a **stream**:
 - ❖ an ordered **sequence** of **bytes** (similar to an array)
 - ❖ can be read, written, and moved through
 - ❖ separate streams to several files can be opened at one time

Opening a File

- ❖ A file can be **opened** for **reading** or **writing**.
- ❖ The filename is associated with a **file pointer**.
- ❖ The pointer points to a structure that contains the file information which is managed by the OS (bookkeeping information that the OS has to share with program so that it can access the file).

Declaring a File Pointer

- ❖ A file pointer is declared using the `FILE` type:


```
FILE *fp;
```

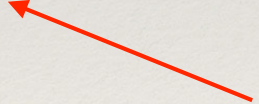
```
FILE *infile;
```

- ❖ The file is opened using `open`.

```
fp = fopen ( "filename", "r" );
```


file pointer


name of the file that you wish to open


mode

File Modes

- ❖ Mode indicates what you can do with the file

Mode	Meaning
r	open existing file for reading ; start at the <i>beginning</i>
w	create a new file for writing
a	append - open an existing file; write at the <i>end</i> of the file
r+	open existing file for reading and writing ; start at the <i>beginning</i>
w+	create new file for reading and writing ; start at the <i>beginning</i> ; truncates existing file
a+	open existing file for reading and writing ; write to the <i>end</i> of the file

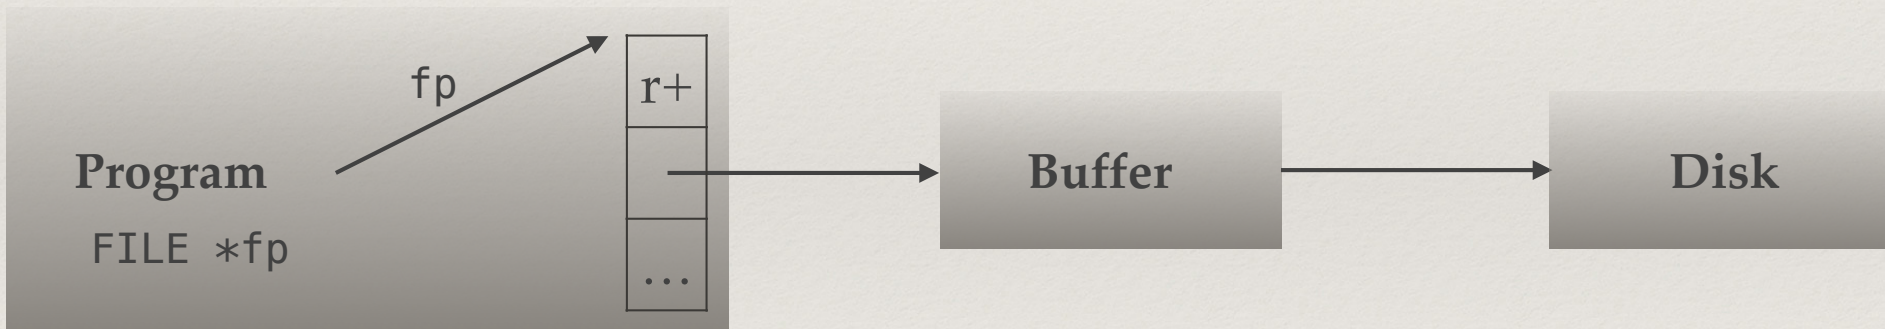
Files, Buffers & C

- ❖ There is not a direct path to the disk.
- ❖ It would be inefficient to actually write to the disk (hardware) after each program write.



Files, Buffers & C

- ❖ There is not a direct path to the disk.
- ❖ It would be inefficient to actually write to the disk (hardware) after each program write.



Reading from a File

- ❖ Reading from a file containing text can be done using the same commands that are used to read from the keyboard (*i.e.* the `stdin` stream).
- ❖ You need to use the “file” versions: `fscanf`, `fgets`.
- ❖ Both commands are exactly the same as reading from the keyboard / `stdin` with the addition of a file pointer.

fscanf and fgets

- ❖ Read individual values from a specified stream (file).

```
fscanf ( fp, "%d %c", &num, &ch );
```

Read from the stream
associated with fp in
the fopen command

Exactly the same as scanf

- ❖ Read a string from an input stream.

```
fgets ( str, 100, fp );
```

char array to
store input

maximum
length of input


read from file
associated with
fp in fopen

Writing to a File

- ❖ Writing text to a file is performed using `fprintf` and `fputs`.
- ❖ Similar to output using `printf` but writes to a file instead of the screen (`stdout`).

```
fprintf ( fp, “%s %d”, stringA, numB );
```

File pointer
from `fopen`



Same as `printf`




```
fgets( strPtr, length, fp );
```


Pointer to string




maximum
length of input



File pointer
from `fopen`



Writing to a File



- ❖ Note: the program could have written to the buffer, but the buffer not yet stored on the disk
- ❖ Buffers are written to disk on
 - `fprintf (fp, “\n”)`
 - `fflush (fp)`

Closing a File

- ❖ When you have finished reading or writing to a file it must be closed using `fclose`.

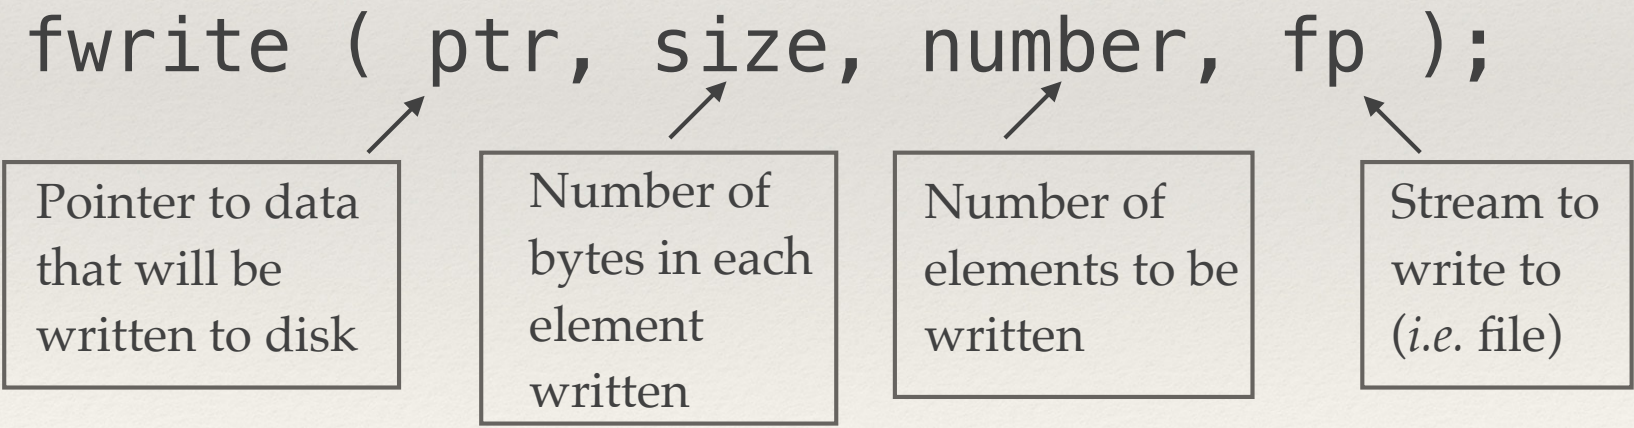
`fclose (fp);`

- ❖ This forces any buffered data to be written to the disk.
- ❖ If the file is not closed then data may be lost (*i.e.* not written to the file).
- ❖ It also frees the memory used to store file information.

Writing Binary Data to a File

- ❖ Writing binary data to a file requires the `fwrite()` command.
- ❖ This writes the data in the form that it is stored in by the system (**raw** bits) and not as ASCII.

```
fwrite ( ptr, size, number, fp );
```



Pointer to data
that will be
written to disk

Number of
bytes in each
element
written

Number of
elements to be
written

Stream to
write to
(*i.e.* file)

Binary File Write Example

```
#include <stdio.h>
int main ( )
{
    int numArray[100], i;
    FILE *fptr;
```

```
fptr = fopen ( "tmpfile", "wb" );
```

open for writing

```
for ( i=0; i<100; i++ ) {
    numArray[i] = i;
}
```

*write 100 elements that are
the size of an int (4 bytes),
taken from array numArray
and put into the file associated
with the file pointer fptr*

```
fwrite ( numArray, sizeof(int), 100, fptr );
```

```
fclose ( fptr );
```

```
}
```

- Creates a 400 byte file:

```
$ ls -l tmpfile
```

```
•-rw-r--r--  1 dastacey  staff  400  6 Dec 14:51 tmpfile
```



```
$ od -d tmpfile
```

0000000	0	0	1	0	2	0	3	0
0000020	4	0	5	0	6	0	7	0
0000040	8	0	9	0	10	0	11	0
0000060	12	0	13	0	14	0	15	0
0000100	16	0	17	0	18	0	19	0
0000120	20	0	21	0	22	0	23	0
0000140	24	0	25	0	26	0	27	0
0000160	28	0	29	0	30	0	31	0
0000200	32	0	33	0	34	0	35	0
0000220	36	0	37	0	38	0	39	0
0000240	40	0	41	0	42	0	43	0
0000260	44	0	45	0	46	0	47	0
0000300	48	0	49	0	50	0	51	0
0000320	52	0	53	0	54	0	55	0
0000340	56	0	57	0	58	0	59	0
0000360	60	0	61	0	62	0	63	0
0000400	64	0	65	0	66	0	67	0
0000420	68	0	69	0	70	0	71	0
0000440	72	0	73	0	74	0	75	0
0000460	76	0	77	0	78	0	79	0
0000500	80	0	81	0	82	0	83	0
0000520	84	0	85	0	86	0	87	0
0000540	88	0	89	0	90	0	91	0
0000560	92	0	93	0	94	0	95	0
0000600	96	0	97	0	98	0	99	0
0000620								

Reading a Binary File

- ❖ Reading binary data from a file is done with `fread()`.
- ❖ `fread()` has the same parameters as `fwrite()`.

`fread (ptr, size, number, fptr);`

array that the
data will be
stored in or
address of a
data structure

size of the
array
elements

number of
array
elements

file to be
read from

Binary File Read Example


```
#include <stdio.h>
int main ( )
{
    int numArray[100];
    FILE *fptr;
```

*given the file created in the previous
program to write a binary file*



```
    fptr = fopen ( "tmpfile", "rb" );
    fread ( numArray, sizeof(int), 100, fptr );
    fclose ( fptr );
```

*reads 100 integers from the
file (tmpfile) into the
array (numArray)*



```
    printf ( "numbers: %d %d\n", numArray[0], numArray[99] );
}
```

```
$ ./binaryFileRead
numbers: 0 99
```

Reading a Binary File

- ❖ The system is really reading / writing the number of bytes equal to `size * number`.
- ❖ It does not read one element at a time.
- ❖ It reads all the requested bytes and places them in the data structure (*e.g.* array).
 - ❖ Boundaries between elements are ignored.
- ❖ In the previous two examples, the file created is a collection of 100 integers (400 bytes).

Reading and Writing Structures

- ❖ Reading and writing structures can be achieved using fread / fwrite.

- ❖ Example: write and read an array of 10 structures to disk

```
#include <stdio.h>
```

```
struct data {  
    char letter;  
    int number;  
    double bignum;  
};
```

```
typedef struct data datatype;
```

```
int main ( )  
{
```

```
    /* code that uses datatype */
```

```
}
```

*replaces struct data with
one word (datatype)*



Reading and Writing Structures

❖ Read

fread

❖ Exam

disk

```
int main ( )
{
    struct data *mydata

    /* code that uses mydata struct pointer */
}
```

same as

```
int main ( )
{
    datatype *mydata

    /* code that uses mydata struct pointer */
}
```

data with
)

```
int main ( )
```

```
{
```

```
    /* code that uses datatype */
```

```
}
```



```
FILE *fptr;  
int i;  
datatype labels[10];
```

```
...
```

```
/*  
 *   Write the data structures  
 */  
fptr = fopen ( "outfile", "wb" );  
fwrite ( labels, sizeof ( datatype ), 10, fptr );  
fclose ( fptr );
```

```
...
```

```
/*  
 *   Read the data structures  
 */  
fptr = fopen ( "outfile", "rb" );  
fread ( labels, sizeof ( datatype ), 10, fptr );  
fclose ( fptr );
```

Reading and Writing Structures

❖ How big is the file `outfile` when written?

❖ `filesize = sizeof (datatype) * 10`

❖ `datatype` contains:

1 char	1 byte
1 integer	4 bytes
1 double	8 bytes
	<hr/>
	13 bytes

❖ The file should be $13 * 10 = 130$ bytes.

❖ But structures can be padded so that each element begins on a 4 byte boundary.

Reading and Writing Structures

❖ Padding

1 char	4 bytes
1 integer	4 bytes
1 double	8 bytes
	16 bytes

❖ Why?

- ❖ Variables that start on a 4-byte boundary are faster to manipulate by the machine.
- ❖ The actual file size is $16 * 100 = 160$ bytes

Reading and Writing Structures

- ❖ The difference between passing a structure versus passing an array to `fread()` and `fwrite()`?
- ❖ Arrays provide their own address
- ❖ Their name = their address

```
datatype a[10];
```

```
datatype b;
```

array name = its address

```
fwrite ( a, sizeof ( datatype ), 10, fptr );
```

```
fread ( &b, sizeof ( datatype ), 1, fptr );
```

must take address of structure (unless supplying a struct pointer)

Sequential Access Files

- ❖ All the files that we have examined have been read / written in a sequence.

- ❖ Sequential Access File

item 1	item 2	item 3	item 4	...	item m
--------	--------	--------	--------	-----	----------


- ❖ item could be a primitive type (int, float, *etc.*) or a structure (struct)
- ❖ records (or items) are read / written in order

Random Access Files

- ❖ But often you do not want to read or write only in order; your algorithm might need to “*jump*” around the file getting records in **random** order.
- ❖ Random Access allows files to be read / written in **any** order.
- ❖ The files themselves are the same (sequence of bytes) as all other files but there are two functions used to control where access occurs.

Random Access Files

- ❖ `fseek()` - move to a location in the file
- ❖ `ftell()` - report to your program its current location in the file
- ❖ Both `fseek()` and `ftell()` use a current position value to indicate where the file will next be accessed.



Current Position

- next read or write will happen at record X

fseek()

fseek (fptr, offset, from);

file pointer

how far to move
in the file

one of 3 values that
tell where to seek
from

fseek() returns 0 on **success**
and non-zero on *failure*

SEEK_SET

start of the file

SEEK_CUR

from the current position

SEEK_END

end of the file

fseek()

```
fseek ( fptr, sizeof(datatype)*3, SEEK_SET );
```

in file
associated
with fptr

move 3 times the
size of the
datatype structure

move from the start
of the file

SEEK_SET

start of the file

SEEK_CUR

from the current position

SEEK_END

end of the file

fseek()

```
fseek (fptr, sizeof(datatype)*-3, SEEK_END);
```

in file
associated
with fptr

move **back** 3 times
the size of the
datatype structure

move from the end
of the file

SEEK_SET

start of the file

SEEK_CUR

from the current position


SEEK_END

end of the file

ftell()

```
long tell ( fptr );
```

- ❖ Return value is the number of bytes from the start of the file
- ❖ Example: in a file of integers (4 bytes each)



- ❖ `ftell()` would return a value of 12


```
#include <stdio.h>
#include <string.h>
```

```
int main ()
{
    struct info {
        char name[10];
        int id;
    };
    typedef struct info record;
    record item;
    record newItem[10];

    FILE *fptr;
    int i;
    char names[40] = "FredJohnMaryFordGertHammWillBillPaulSamm";
```

```
fptr = fopen ( "data", "w+" );
```

*Open new file for
writing and reading*

```
for ( i=0; i<10; i++ ) {
    strncpy ( item.name, &names[i*4], 4 );
    strncpy ( &item.name[4], "\0", 1 );
    item.id = i;
```

```
    fwrite ( &item, sizeof(item), 1, fptr );
```

Write a record to the file

```
}
```

```
fseek ( fptr, 0, SEEK_SET );
```

Move to the beginning of the file

```
for ( i=0; i<10; i++ ) {
```

```
    fread ( &newItem[i], sizeof(record), 1, fptr );
```

Read a record from the file

```
}
```

```
for ( i=0; i<10; i++ ) {
```

```
    printf ( "(%d) %s - %d\n", i+1, newItem[i].name, newItem[i].id );
```

```
}
```

*Start at the
beginning of
the new file
and write
10 records*

*Move to the end of Record 4
(starting from the beginning
of the file)*

A `fseek (fptr, (sizeof(item) * 4), SEEK_SET);`

Read Record 5

B `fread (&item, sizeof(item), 1, fptr);`

`printf ("Record 5: %s - %d\n", item.name, item.id);`

*Move 3 records backwards
from the current location*

C `fseek (fptr, sizeof(item) * -3, SEEK_CUR);`

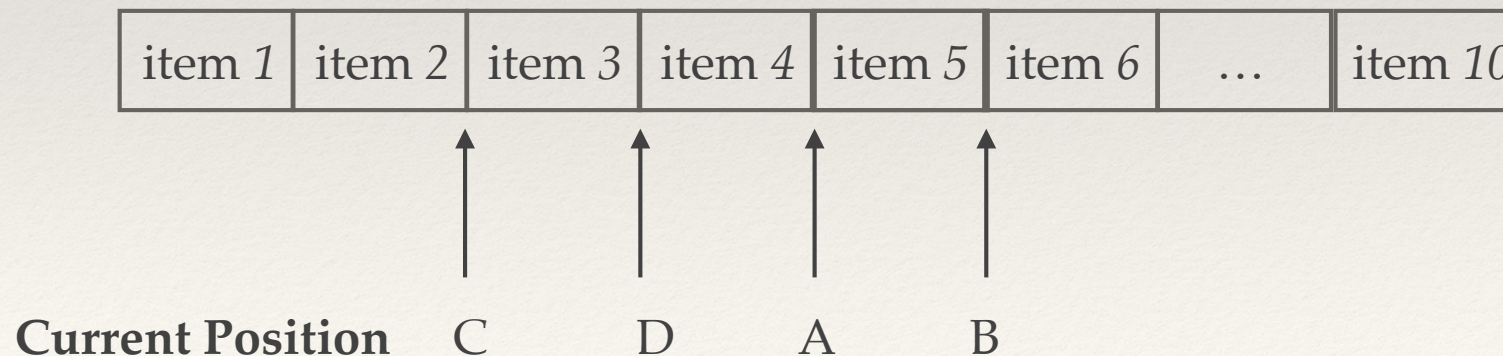
Read Record 3

D `fread (&item, sizeof(item), 1, fptr);`

`printf ("Record 3: %s - %d\n", item.name, item.id);`

`fclose (fptr);`

}




```
$. /seekTell
```

```
(1) Fred - 0
```

```
(2) John - 1
```

```
(3) Mary - 2
```

```
(4) Ford - 3
```

```
(5) Gert - 4
```

```
(6) Hamm - 5
```

```
(7) Will - 6
```

```
(8) Bill - 7
```

```
(9) Paul - 8
```

```
(10) Sam - 9
```

```
Record 5: Gert - 4
```

```
Record 3: Mary - 2
```

Note: The file is padded since the structure was not a multiple of 4 bytes.

1	2	3	4	5	6	7	8	9	10	P	P	1	2	3	4
---	---	---	---	---	---	---	---	---	----	---	---	---	---	---	---

name _____ *Padding* — **id** _____